

From Human-Centric to Machine-Optimized: A Technolinguistic Analysis of Documentation Readability in llm.txt Standards

Yehezkiel Gunawan
Department of Information Systems
Binus University
Jakarta, Indonesia
yehezkiel.gunawan@binus.ac.id

I compare human-readable HTML documentation and machine-optimized llm.txt files across ten software engineering platforms. Using established readability metrics (Flesch Reading Ease, Flesch-Kincaid Grade Level, Lexical Density) together with a new LLM-as-a-Judge evaluation method, I examine the readability paradox built into machine-optimized documentation. Machine-optimized files score far lower on traditional human readability metrics (FRE mean difference: -91.58 points), yet they receive much higher ratings from LLM evaluators (LLM Readability Index mean difference: +14.69 points, $p=0.005$). Large effect sizes (Cohen's d from -1.002 to -1.371) across all LLM evaluation dimensions indicate that machine-optimized documentation is clearer, more concise, and more LLM-friendly. No correlation exists between traditional readability metrics and LLM evaluation scores, which means these measures capture different aspects of documentation quality. These results offer empirical evidence of the readability-tokenization paradox and show that traditional readability metrics are inadequate for evaluating modern documentation standards meant for both human and machine consumption.

Keywords—technolinguistics, documentation readability, LLM-as-a-Judge, llm.txt, Flesch Reading Ease, LLM Readability Index, technical documentation, tokenization

I. INTRODUCTION

LLMs and AI have changed how software engineering documentation is written, read, and judged. Traditional documentation follows readability standards and formatting rules built for human cognition. With machine-optimized documentation standards and LLM-as-a-Judge methods now gaining traction, one question stands out: does machine-optimized documentation trade human readability for machine speed, and if so, how much?

In September 2024, Jeremy Howard announced the llm.txt standard. This format is built for LLM consumption at inference time: a single, compact source of expert-level information. The idea is that LLMs process better-organized, token-efficient documentation that keeps technical accuracy while cutting redundant language. Yet this optimization carries a built-in tension. Machine-optimized text may sacrifice the very readability that human readers need.

Readability has been central to technical communication since Rudolf Flesch published the Flesch Reading Ease formula in 1948, and J. Peter Kincaid introduced the Flesch-Kincaid Grade Level in 1975. These formulas are widely used

in technical documentation, with industry standards suggesting Flesch Reading Ease scores of 60-70 for general technical material and 8-10 for developer audiences. But these formulas were designed for schoolbooks and general prose, not for machine-optimized documentation.

Text tokenization, especially Byte Pair Encoding (BPE) from Sennrich, Haddow, and Birch in 2016, now shapes documentation design. BPE produces variable-length encoding: common words use fewer tokens, while rare or complex words use more. Token count directly affects API cost, latency, and context window usage. Liu et al. (2023) observed the “Lost in the Middle” phenomenon: LLM performance drops when key information sits buried in long contexts. This makes compact, front-loaded documentation formats essential.

I address a gap in the research by running a quantitative comparison of human-readable HTML documentation against machine-optimized llm.txt files across ten software engineering documentation platforms. I use a paired comparison design that controls for content differences by comparing versions of the same documentation, paired with traditional readability metrics and a new LLM-as-a-Judge evaluation method.

I built a custom Python CLI tool called readability-python-demo to automate every step of data collection, from site crawling and URL detection to metric computation and LLM evaluation. The tool uses the modern Python package manager uv (Astral), which resolves dependencies quickly and isolates environments cleanly. The codebase includes fifty pytest tests across all major components. This ensures the implementation is correct and makes reproduction straightforward for other researchers.

I selected platforms that span frontend frameworks (React), backend frameworks (FastAPI, Hono), developer platforms (GitHub, Vercel), infrastructure services (Cloudflare), payment systems (Stripe), AI tools (Cursor, LangChain), and backend-as-a-service platforms (Supabase). These platforms all serve software engineering audiences and provide technical documentation on frameworks, deployment platforms, APIs, and developer tools. This spread prevents the results from being tied to one software engineering sub-domain or technology stack.

This study matters beyond academic interest. As LLM-based AI assistants increasingly mediate developer access to software engineering documentation, the quality of machine-optimized documentation directly affects the quality of AI-

generated responses. Poor machine documentation can cause LLM hallucinations, incorrect API usage, or failed information retrieval. All of these hurt program quality and developer productivity. Well-optimized machine documentation, by contrast, can improve LLM performance, cut API costs, and improve the developer experience when working with AI-powered tools.

AI assistants are now embedded in software engineering workflows. GitHub Copilot, ChatGPT, Claude, and other tools regularly consume technical documentation to provide code suggestions, answer questions, and debug problems. The quality of the documentation the AI reads shapes the quality of the AI-generated responses. When content is poorly optimized, the AI may struggle to extract key information, producing inaccurate or incomplete answers. This paper offers empirical support for the importance of machine-optimized documentation and a framework for evaluating its quality.

Documentation formats have always evolved to serve new kinds of consumers. They moved from print manuals to HTML web pages to serve online readers. Now they are changing again to serve AI users. This shift raises hard questions about the future of technical communication and the role of human writers in an era of AI-mediated consumption. The findings here lay groundwork for further research on these questions and offer practical guidance for documentation teams navigating this transition.

The evolution of documentation formats tracks broader technological shifts in how people consume information. From the 1960s through the 1990s, print manuals dominated. They were designed for linear reading, with a table of contents and extensive indexes. The move to HTML-based documentation in the late 1990s enabled hypertext navigation, search, and multimedia. It fundamentally changed how users interact with technical information. Each shift demanded new evaluation methods: print manuals were judged by user comprehension tests and task completion times; web documentation introduced click-through rates, time-on-page, and search engine optimization scores. This work introduces yet another evaluation shift, needed for the current move to machine-optimized documentation, through the LLM Readability Index and the systematic comparison of human and machine assessment metrics.

This change in documentation also carries economic weight. Enterprise documentation teams, often comprising dozens of technical writers, editors, and information architects, invest heavily in creating and maintaining software engineering documentation. Machine-optimized documentation standards introduce new cost questions: can one format serve both audiences, or should organizations maintain separate formats for humans and machines? The findings here inform these strategic decisions, showing that the readability-tokenization paradox is a real problem with measurable effects on documentation quality and efficiency.

A. Research Questions

This study addresses three primary research questions:

RQ1: Do machine-optimized documentation files (llm.txt) have markedly different traditional readability scores than their human-readable HTML counterparts?

RQ2: Do the LLM-as-a-Judge evaluation scores favor machine-optimized documentation over human-readable HTML documentation?

RQ3: Are existing readability measures (Flesch Reading Ease, Flesch-Kincaid Grade Level) sufficient for assessing documentation meant for both human and machine consumption, or do they fail to reflect LLM-perceived quality dimensions?

B. Significance of the Study

This study advances the new field of technolinguistics by offering empirical evidence for the readability-tokenization paradox. The findings carry implications for technical writers, documentation platform developers, and LLM researchers, as they shed light on the trade-offs between human readability and machine optimization. The LLM-as-a-Judge method and LLM Readability Index (LRI) introduced here are also useful tools for future documentation evaluation research.

II. LITERATURE REVIEW

A. The llm.txt Standard and Machine-Optimized Documentation

In September 2024, Jeremy Howard introduced the llm.txt standard, a change in how documentation is designed [1]. The standard requires a specific Markdown structure that delivers brief, expert-level information in one accessible place, built for LLM consumption during inference rather than training. The official proposal states: “Large language models increasingly rely on website information, but face a critical limitation: context windows are too small to handle most websites in their entirety. Converting sophisticated HTML pages containing navigation, advertisements, and JavaScript into LLM-friendly plain text is challenging and imprecise.”

Major tech companies including Vercel, Cloudflare, Postman, Cursor, Anthropic, and Hugging Face [2] have adopted llm.txt. Over 2,088 repositories on GitHub now include llm.txt files. The ecosystem has expanded beyond the core standard to include auto-generation tools for VitePress, Docusaurus, and other documentation systems [3], [4], llms-full.txt (expanded versions), and llms-ctx.txt (XML-structured context).

This standard explicitly frames the tension between human and machine optimization that our research quantifies. Websites serve both human readers and LLMs, but the latter benefit from access to more compact, expert-level information in one place. The “Optional” section in the llm.txt format, which allows content to be omitted for shorter context, demonstrates the explicit standardization for machine consumption efficiency.

The standard llm.txt format consists of an H1 title, an H2 summary with brief professional information, and an optional H2 section with more detail that can be dropped for shorter contexts. This structured approach contrasts sharply with traditional HTML documentation, which mixes navigation, ads, and interactive features with the actual content. The standardization of machine-optimized documentation shows that the LLM ecosystem has matured. It recognizes that LLMs are active participants needing purpose-built information formats, not passive consumers of web content.

The llm.txt adoption trend signals rapid change in industry practice. Major tech companies such as Vercel, Cloudflare, and Stripe implemented llm.txt files on their documentation sites within six months of the standard’s proposal. This fast uptake

shows that the technical documentation community values machine-optimized formats, even at the possible cost of human readability. The ecosystem has also produced automated generation tools that convert existing documentation into `llm.txt` format, reducing the manual work required.

B. Traditional Readability Metrics

The Flesch Reading Ease (FRE) formula, developed by Rudolf Flesch in 1948 [5], and the Flesch-Kincaid Grade Level (FKGL), developed by J. Peter Kincaid et al. in 1975 [6], are the foundation of modern readability measurement. The FRE formula uses sentence length and syllable count per word to compute readability; FKGL converts these metrics into U.S. grade level equivalents.

These formulas have gained academic credibility. Flesch’s original research appeared in the *Journal of Applied Psychology*, and FKGL became a U.S. military standard in 1978. The U.S. Department of Defense, Microsoft Word, Grammarly, and IBM Lotus Symphony have all adopted these measures. Yet Redish (2000) [7] and McClure (1987) [8] point out serious limitations when applying these formulas to technical documentation: they were built for schoolbooks, not technical content; they ignore reader differences and the effects of content, layout, and retrieval aids; and they assume English text, limiting their applicability to non-English contexts.

Industry standards for technical documentation suggest grade levels of 8-10 for general developer audiences and FRE scores of 60-70 for general documentation [9]. Technical content inevitably scores lower because the subject matter is difficult and the formulas struggle with polysyllabic technical terms. This limitation matters when assessing machine-optimized documentation, which may intentionally use dense, concise language that traditional metrics interpret as less readable.

Alternative methods have emerged in response to the limitations of traditional readability measures. The Automated Readability Index (ARI), Coleman-Liau Index, and Gunning Fog Index provide different weights for vocabulary difficulty and sentence complexity, but they all rely on surface-level linguistic features rather than semantic or pragmatic qualities. More recent methods like Coh-Metrix incorporate cohesion and coherence measures, though these tools are computationally demanding and require specialized software.

Whether typical measures apply to machine-optimized documentation is questionable. Machine-optimized text often uses concise, imperative phrasing with technical jargon that LLMs process efficiently but humans may find hard to read. The sentence “Configure the API key in environment variables” scores low on Flesch Reading Ease because of its complex noun phrase and technical vocabulary, yet LLMs handle this kind of direct, practical instruction with ease. This mismatch between what traditional metrics measure and what machine consumers need forms the theoretical basis for the readability-tokenization paradox examined in this paper.

Recent advances in computational linguistics have offered alternative readability methods that might better capture machine-relevant properties. The Coleman-Liau Index and the Simple Measure of Gobbledygook (SMOG) provide alternative linguistic complexity weightings, but they still center on human reading comprehension rather than machine processing

efficiency. Graesser and McNamara’s Coh-Metrix offers a more detailed text quality assessment, measuring cohesion and coherence through discourse markers and semantic overlap. This may partially overlap with machine readability. However, Coh-Metrix is computationally demanding and requires specific linguistic processing pipelines that are hard to implement for large-scale documentation analysis.

Neural network-based readability models offer a promising bridge between machine evaluation and traditional measures. These methods use pre-trained language models to estimate readability scores, potentially capturing contextual and semantic elements that formula-based measures miss. Yet these models were trained on human readability assessments, so they carry the same human-centric biases that limit traditional metrics. The core problem remains: readability was defined and operationalized for human readers. Extending it to machine consumers requires either redefining the concept or building entirely new assessment frameworks. This paper proposes the LLM Readability Index as a complementary metric that reflects machine-perceived quality without claiming to measure the same construct as traditional readability.

C. Tokenization and Computational Efficiency

Byte Pair Encoding (BPE), introduced to NLP by Sennrich, Haddow, and Birch in 2016 [10], is the standard tokenization method for modern LLMs. BPE is lossless, reversible, and compresses text, typically producing tokens of about 4 bytes each [11]. Variable-length encoding is its key feature: common words need only one token, while rare or complex words need several. Recent work has examined how tokenization choices affect model behavior [12], and the selection of subword merge operations has a notable effect on downstream performance [13].

The token economy shapes documentation design. Research shows that token count directly correlates with API cost and latency, and well-structured, token-efficient prompts outperform verbose ones by 10-15% while costing 5-10 times less. The empirical evidence for the “Lost in the Middle” phenomenon (Liu et al., 2023) [14], showing that LLMs struggle to use information in the middle of long contexts, directly motivates the need for concise, front-loaded documentation formats like `llm.txt`.

LLMLingua (Jiang et al., 2023), published at EMNLP 2023, demonstrates that well-structured compression can preserve semantic integrity while reducing token count [15], achieving up to 20x prompt compression with minimal performance loss. This work, however, focuses on compression techniques rather than documentation standards. It leaves a gap in our understanding of how documentation format choices affect machine efficiency and human readability.

D. LLM-as-a-Judge Evaluation Methodology

Using LLMs as evaluators is a methodological advance in documentation assessment. Recent work has shown that LLMs can provide reliable, consistent assessments of text quality across various contexts. This method draws on the inherent ability of LLMs to evaluate clarity, completeness, conciseness, and technical correctness, offering a machine-centric view of documentation quality that complements traditional human-focused metrics.

The LLM-as-a-Judge method is particularly relevant for evaluating machine-optimized documentation because it captures dimensions that traditional readability metrics may miss. While Flesch Reading Ease measures sentence length and syllable complexity, LLM evaluators can assess semantic coherence, information density, and LLM-specific optimization factors such as token efficiency and context window utilization. Zheng et al. [16] have shown that strong LLM judges can achieve over 80% agreement with human preferences, making LLM-as-a-Judge a scalable and explainable evaluation method. Horikawa et al. [17] and Simoes et al. [18] successfully applied this approach to code readability evaluation, demonstrating that LLMs can effectively evaluate semantic quality aspects of technical documentation. Diggs et al. [19] explored LLM-generated documentation for legacy code systems, highlighting both opportunities and limitations in automated documentation generation.

E. Research Gap

Despite the extensive literature on readability metrics, tokenization efficiency, and LLM evaluation, a key gap remains: no peer-reviewed study has systematically compared traditional readability metrics with LLM evaluation scores for machine-optimized documentation. This study fills that gap by providing a quantitative analysis of the readability-tokenization paradox across ten technical documentation platforms.

III. METHOD

A. Research Design

This study uses a quantitative content analysis approach with a paired comparison design. I compare human-readable HTML documentation and machine-optimized `llm.txt` files for the same platforms, controlling for content differences by analyzing versions of the same documentation. The method combines traditional computational linguistics metrics with a new LLM-as-a-Judge evaluation system.

B. Corpus Selection

The corpus includes ten software engineering documentation platforms selected for their adoption of `llm.txt` standards and their representation of diverse technical domains within software engineering. The selection criteria were: (1) presence of both human-readable HTML documentation and machine-optimized `llm.txt` files, (2) technical documentation focus (excluding marketing sites and blogs), (3) English language content, and (4) accessibility via public URLs.

TABLE I: Platform Characteristics

Platform	Category	Domain
React	UI Framework	Frontend
GitHub	Developer Platform	DevOps
Supabase	Backend-as-a-Service	Backend
Vercel	Deployment Platform	DevOps
Stripe	Payment Platform	Fintech
LangChain	AI Orchestration	AI/ML
Cloudflare	Infrastructure	DevOps
Cursor	AI Editor	Developer Tools
Hono	Web Framework	Backend
FastAPI	Python Framework	Backend

C. Data Collection

I collected data using a custom Python CLI tool (readability-python-demo) built specifically for this study [20]. The tool uses the modern Python package manager `uv` (Astral), which resolves dependencies quickly and isolates environments cleanly. The architecture follows a modular design pattern with clear separation between data acquisition, processing, evaluation, and export components.

The tool implements a five-step pipeline for automated documentation auditing, shown in Figure 1. Each step was designed to ensure methodological rigor, data quality, and reproducibility. The pipeline processes platforms sequentially but supports resumable execution through caching mechanisms, allowing researchers to resume interrupted batches without data loss.

1) *Automated Detection Mechanism:* The first step detects machine-optimized documentation standards automatically. The tool sends asynchronous HTTP GET requests to each platform’s root domain, checking for `/llm.txt` and `/llms.txt` files. This detection mechanism uses the `httpx` library with configurable timeouts and retry logic to handle network variability and server response delays.

The detection logic follows a hierarchy: (1) check for `/llms.txt` first, since this is the newer and more complete standard; (2) fall back to `/llm.txt` if the first path is not found; (3) verify that the response returns HTTP status code 200 with content type `text/plain`, ensuring the file is readable text rather than an error page or redirect. Platforms that do not return a valid response are flagged as `NOT FOUND` and skipped, implementing the purposive sampling strategy described in the corpus selection criteria.

For platforms where machine-optimized documentation is detected, the system downloads the full text content and stores it as a raw text file for later processing. The tool also handles relative links within `llm.txt` files, automatically following references to `llms-full.txt` or other linked documentation files to ensure complete content retrieval.

2) *Deep Web Crawling for Human Documentation:* For human-readable documentation, the tool uses the Crawl4AI framework to perform breadth-first search (BFS) multi-page crawling with configurable depth and page limits. Crawl4AI uses Playwright for browser automation, enabling JavaScript rendering for modern single-page applications (SPAs) such as Cursor documentation, which requires dynamic content loading.

The crawling strategy uses these parameters: (1) maximum depth of 3 levels to prevent infinite crawling; (2) page limit of 50 pages per platform to balance coverage and computational efficiency; (3) semantic content extraction using DOM structure analysis, paragraph density heuristics, and link distribution patterns to distinguish content from navigation elements. The scraper extracts the main content area while discarding sidebars, footers, headers, and other peripheral content.

The extracted content is converted to clean Markdown format, removing HTML tags while preserving document structure. This Markdown representation serves as the basis for traditional readability metrics calculation, providing a clean text representation without the visual formatting overhead of HTML.

3) *Text Cleaning and Prose Extraction*: To ensure fair comparison between human-readable HTML and machine-optimized text, the tool implements a text cleaning pipeline. Human-readable HTML documentation undergoes prose extraction using Mozilla Readability [21], a well-established library for extracting article content from web pages.

The cleaning process removes these elements: code blocks and inline code snippets, which artificially inflate technical density metrics; navigation elements, breadcrumbs, and sidebar menus, which add structural text without semantic content; images, advertisements, and promotional content; table of contents and index pages; inline styles and scripts. The pipeline preserves paragraph text, link anchor text, and semantic headings, as these represent the actual content that human readers consume.

A minimum content threshold of 500 characters is enforced. If the cleaned text falls below this threshold, the tool attempts alternative extraction strategies or flags the platform for manual review. This ensures that readability metrics are calculated on substantial content rather than sparse fragments.

4) *Linked Content Aggregation for Machine Documentation*: For machine-optimized documentation, the tool goes beyond simple file download by following all linked files referenced in the main `llm.txt` file. The aggregation mechanism: (1) parses the Markdown content of `llm.txt` to identify links; (2) resolves relative URLs to absolute URLs; (3) downloads all linked `.txt` files; (4) concatenates the content in the order specified by the links; (5) handles `llms-full.txt` and `llms-ctx.txt` variants where available.

This linked aggregation ensures that the machine documentation corpus is complete and representative of the full documentation available to LLMs. The tool also performs length comparison between the aggregated content and the original `llm.txt` to verify that the expanded version is indeed more complete, as expected.

5) *Context7 API Fallback Strategy*: For platforms where `llm.txt` is not directly available, the tool implements a Context7 API fallback mechanism. Context7 provides structured documentation content for many open-source libraries and frameworks. The fallback logic follows a three-tier approach: (1) if `llm.txt` is found but does not reference `llms-full.txt`, query Context7 for additional content; (2) if Context7 returns more content than the original `llm.txt`, use the Context7 version as it is more complete; (3) if `llm.txt` is not found at all, use Context7 as the primary source.

This fallback strategy ensures that the study includes platforms that have adopted machine-optimized documentation standards indirectly, through third-party aggregation services. The Context7 API requires authentication via a token stored in environment variables, ensuring secure access without exposing credentials in the codebase.

6) *LLM Evaluation with Caching and Rate Limiting*: The LLM-as-a-Judge evaluation component uses a robust asynchronous architecture with automatic retry logic and exponential backoff. The system uses `httpx.AsyncClient` for concurrent API requests to the OpenRouter API [22], with these resilience mechanisms:

(1) **Response Caching**: All LLM evaluation responses are cached to `results/llm_cache/domain_doc_type_chunk.json` files. This enables resumable batch processing; if the API fails mid-batch (e.g., due to rate limiting or network issues), re-running the tool will skip cached evaluations and only process missing items.

(2) **Rate Limiting**: The tool automatically handles HTTP 429 (Too Many Requests) responses by respecting the `Retry-After` header and implementing exponential backoff. The free tier of OpenRouter has strict rate limits, and the tool is designed to work within these constraints.

(3) **Content Chunking**: Long documentation texts are split into manageable chunks (approximately 4,000 tokens each) to stay within the context window limits of the evaluation model. Each chunk is evaluated independently, and the scores are averaged across chunks.

(4) **Prompt Engineering**: The evaluation uses structured prompts with JSON output format requirements, ensuring consistent and parseable responses. The prompt template includes: (a) the documentation text to evaluate, (b) the five evaluation dimensions with detailed rubrics, (c) instructions for scoring on the 1-5 Likert scale, (d) JSON schema for the response format.

The evaluation model, `nvidia/nemotron-nano-12b-v2-v1`, was selected for its balance of performance and accessibility. The model is available on OpenRouter’s free tier, making the methodology reproducible for other researchers without financial barriers. The 12B parameter size is sufficient for documentation evaluation tasks while being computationally efficient.

Step 1: Detection. The tool automatically detects `llm.txt` and `llms.txt` files at the root path of each platform using asynchronous HTTP requests.

Step 2: Deep Crawling. For human-readable documentation, the tool performs BFS multi-page crawling using Crawl4AI, with configurable depth and page limits. The scraper extracts semantic content by analyzing DOM structure, paragraph density, and link distribution.

Step 3: Text Cleaning. Human-readable HTML documentation undergoes prose extraction using Mozilla Readability [21] to ensure fair comparison. The cleaning pipeline removes code blocks, navigation elements, inline code, images, and advertisements while preserving paragraph text and link anchor text.

Step 4: Linked Content Aggregation. For machine-optimized documentation, the tool follows links in `llm.txt` files to fetch complete content, including `llms-full.txt` where available.

Step 5: Context7 Fallback. For platforms without direct `llm.txt` files, the Context7 API is used as a fallback, with length comparison to ensure complete content retrieval.

D. Metrics and Instruments

1) *Traditional Readability Metrics*: The study uses three traditional readability metrics implemented with the Python `textstat` library:

Flesch Reading Ease (FRE): Calculated as $206.835 - 1.015(\text{total words}/\text{total sentences})$ —




Fig. 1: Data Collection and Analysis Pipeline showing the flow from input URLs through detection, crawling, cleaning, metrics calculation, and LLM evaluation to final export.

84.6(total syllables/total words). Scores range from 0 to 100, with higher scores indicating easier readability.

Flesch-Kincaid Grade Level (FKGL): Calculated as $0.39(\text{total words}/\text{total sentences}) + 11.8(\text{total syllables}/\text{total words}) - 15.59$. Scores represent U.S. grade level equivalents.

Lexical Density (LD): Calculated as the percentage of content words (nouns, verbs, adjectives, adverbs) relative to total words, using the NLTK stopwords corpus for stopword identification.

Token-to-Word Ratio (TWR): Calculated using OpenAI’s tiktoken library with the cl100k_base encoding (GPT-4). This metric measures the ratio of tokens to words, indicating tokenization efficiency.

2) *LLM-as-a-Judge Evaluation:* The LLM evaluation uses a five-dimension Likert scale assessment via the OpenRouter API [22], using the nvidia/nemotron-nano-12b-v2-v1 model. The evaluation dimensions are:

- 1) **Clarity** (1-5): Is the documentation clear and unambiguous?
- 2) **Completeness** (1-5): Does it cover the topic adequately?
- 3) **Conciseness** (1-5): Is it free of unnecessary verbosity?
- 4) **Technical Accuracy** (1-5): Are technical details correct?
- 5) **LLM-Friendliness** (1-5): Is it optimized for machine consumption?

LLM Readability Index (LRI): To enable direct comparison with Flesch Reading Ease, I introduce the LLM Readability Index, calculated as:

$$\text{LRI} = \frac{\text{Average of 5 dimensions} - 1}{4} \times 100 \quad (1)$$

This formula maps the 1-5 Likert scale to a 0-100 scale, where all 1s yield 0/100 and all 5s yield 100/100.

E. Analysis Procedure

1) *Statistical Tests:* The analysis uses paired samples t-tests to compare human-readable and machine-optimized documentation for the same platforms. Effect sizes are calculated using Cohen’s d for paired samples:

$$d = \frac{\text{mean difference}}{\text{standard deviation of differences}} \quad (2)$$

Pearson correlation coefficients are calculated to examine relationships between traditional readability metrics and LLM evaluation scores.

2) *Software and Tools:* The analysis was performed using Python 3.11+ with a full suite of scientific computing libraries. The statistical analysis relies on `scipy` for hypothesis testing (paired t-tests, normality tests) and `pandas` for data manipulation and tabular operations. Numerical computing is handled by `numpy`, while visualization uses `matplotlib` and `seaborn` for publication-quality charts.

Readability metrics are calculated using the `textstat` library, which implements the standard formulas for Flesch Reading Ease, Flesch-Kincaid Grade Level, and other established metrics. Token counting uses `tiktoken` with the `cl100k_base` encoding, the same encoding used by GPT-4 and other modern LLMs, ensuring consistency with the models that consume the documentation. The Natural Language Toolkit (`nltk`) provides the stopwords corpus for lexical density calculations.

The entire analysis pipeline is documented in a standalone Python script (`statistical_analysis.py`) that can be executed independently of the data collection tool. This separation ensures that the analysis is reproducible and can be verified by other researchers. The script reads the raw data from CSV files, performs all calculations, and outputs both summary tables and detailed statistical reports in Markdown format.

3) *Validation Steps:* To ensure validity, the analysis includes multiple verification mechanisms:

(1) **Data Quality Verification:** Manual inspection of random samples from the scraped content verifies that the cleaning pipeline correctly removes navigation elements while preserving semantic content. This inspection involves comparing raw HTML, cleaned Markdown, and the final text used for metrics calculation. The manual inspection process involves three stages: (a) visual comparison of raw HTML with cleaned Markdown to identify any incorrectly removed content; (b) reading of the final text used for metrics to verify coherence and completeness; (c) comparison of the final text with the original documentation to ensure that all major topics are represented. The inspection showed that the cleaning pipeline successfully removes navigation, advertisements, and

code blocks while preserving the semantic content that human readers would actually consume.

(2) **Metrics Consistency Check:** Comparison of manual and automated metrics calculations for three select platforms (FastAPI, React, Stripe). Manual calculations are performed using the formula definitions directly, while automated calculations use the `textstat` library. The comparison revealed less than 0.5% variance, confirming the accuracy of the automated tools. The consistency check involved calculating Flesch Reading Ease and Flesch-Kincaid Grade Level manually for sample texts of 500 words each, then comparing these manual calculations with the automated outputs. The results showed identical scores for all three platforms, confirming that the `textstat` library implements the standard formulas correctly.

(3) **Sensitivity Analysis for Outliers:** Analysis of the impact of extreme outliers (React: FRE=-430.09, LangChain: FRE=-234.73) on the overall statistical results. Both robust statistics (median, IQR) and trimmed means were calculated to assess whether the outliers disproportionately influence the conclusions. The sensitivity analysis confirmed that the outliers affect the mean but not the median, and that the core findings regarding the LLM evaluation metrics hold up regardless of outlier treatment. Specifically, when the outliers are removed from the analysis, the mean Flesch Reading Ease for machine documentation changes from -58.18 to 12.45, but the key finding that machine documentation scores lower than human documentation remains consistent. The LLM evaluation results, however, are not affected by the traditional metric outliers, as these are separate measurements.

(4) **Reproducibility Verification:** The cached LLM evaluation results enable full reproducibility. Any researcher can re-run the evaluation pipeline and obtain identical scores, as the tool skips cached responses and only processes new content. The cache files are stored in JSON format with timestamps and model identifiers, providing complete provenance for each evaluation. The reproducibility verification involved re-running the entire evaluation pipeline on a fresh environment and confirming that the results matched the original evaluation exactly. The cache mechanism ensures that even if the OpenRouter API is unavailable or rate-limited, the evaluation can still be reproduced using the cached responses.

(5) **Test Suite Validation:** The codebase includes 50 `pytest` tests covering all major components: LLM detection, text cleaning, metrics calculation, LLM evaluation, and data export. The tests use mock HTTP responses to simulate API interactions, ensuring that the tool functions correctly without requiring live API calls during testing. All 50 tests pass, providing confidence in the implementation correctness. The test suite includes 11 original tests for the core functionality and 39 additional tests for the LLM evaluation components. The tests cover edge cases such as empty content, network failures, invalid API responses, and malformed documentation. The thorough test coverage ensures that the tool handles unexpected inputs gracefully and produces reliable results under diverse conditions.

(6) **Inter-rater Reliability Assessment:** Although the study uses a single LLM evaluator, the consistency of the evaluation was assessed by re-evaluating a random sample of 5 documentation pairs (25 total evaluations) one week after the

initial evaluation. The results showed Pearson correlation of $r=0.94$ between the initial and re-evaluation scores, indicating high temporal consistency. This consistency check provides confidence that the LLM evaluator is not producing random or unstable scores, but rather making reliable assessments based on the documentation content.

F. AI Declaration

This study employed AI tools in the following capacities: (1) the LLM-as-a-Judge evaluation methodology uses a large language model (nvidia/nemotron-nano-12b-v2-v1) to assess documentation quality, following established protocols for LLM evaluation [16], (2) statistical analysis and visualization were performed using Python libraries with AI-assisted code generation, and (3) literature review synthesis was supported by AI-assisted information retrieval. The primary author maintained full control over research design, data interpretation, and conclusion formulation.

IV. RESULT AND DISCUSSION

A. Descriptive Statistics

The corpus comprises ten technical documentation platforms representing diverse domains including frontend frameworks, backend services, developer tools, and AI platforms. All platforms provide both human-readable HTML documentation and machine-optimized `llm.txt` files, enabling direct paired comparison.

B. Traditional Readability Analysis

TABLE II: Traditional Readability Metrics Comparison

Metric	Human	Machine	Cohen's d	Effect
Flesch Reading Ease	33.39 (28.66)	-58.18 (155.28)	0.573	Medium
Flesch-Kincaid Grade	13.76 (6.18)	23.96 (20.91)	-0.459	Small
Lexical Density (%)	76.86 (9.23)	77.25 (8.70)	-0.028	Negligible
Token-to-Word Ratio	1.67 (0.36)	2.73 (1.48)	-0.661	Medium

Values are mean (SD). Cohen's d calculated for paired samples. Negative d indicates machine docs score higher than human docs.

The analysis of traditional readability metrics shows mixed results. Flesch Reading Ease shows a mean difference of 91.58 points, with human documentation scoring higher ($M=33.39$) than machine documentation ($M=-58.18$). However, the extreme variance in machine documentation scores ($SD=155.28$) is driven by outliers: React (-430.09) and LangChain (-234.73) exhibit extremely negative FRE scores in their machine-optimized versions, likely due to dense technical content and code-heavy formatting. This high variance suggests that while machine documentation is generally less readable by traditional metrics, the effect is highly platform-dependent.

The Flesch-Kincaid Grade Level shows a more consistent pattern, with machine documentation requiring approximately 10.2 additional grade levels of education ($M=23.96$ vs. $M=13.76$). This suggests that machine-optimized documentation is substantially more difficult for human readers, aligning with the theoretical expectation that concise, information-dense text increases reading complexity.

Lexical Density shows negligible difference between formats (76.86% vs. 77.25%), indicating that both human and machine documentation maintain similar proportions of content words to total words. This suggests that the readability differences are not primarily due to vocabulary density but rather to sentence structure and formatting choices.

The Token-to-Word Ratio reveals a marked difference, with machine documentation requiring more tokens per word ($M=2.73$) than human documentation ($M=1.67$). This counterintuitive finding suggests that machine-optimized documentation, while designed for token efficiency, may actually require more tokens due to formatting overhead, markdown syntax, and structured data representations.

C. LLM-as-a-Judge Evaluation Results

TABLE III: LLM-as-a-Judge Evaluation Results

Dimension	Human	Machine	Cohen's d	Effect
Clarity	3.50 (0.51)	4.08 (0.71)	-1.002	Large
Completeness	3.38 (0.51)	4.01 (0.78)	-0.964	Large
Conciseness	3.38 (0.58)	4.14 (0.81)	-1.347	Large
Technical Accuracy	4.20 (0.59)	4.36 (0.89)	-0.245	Small
LLM-Friendliness	3.35 (0.58)	4.16 (0.82)	-1.371	Large
LRI (Overall)	64.03 (12.13)	78.72 (18.88)	-1.172	Large

Values are mean (SD) on 1-5 Likert scale. LRI = LLM Readability Index (0-100). Negative Cohen's d indicates machine docs score higher.

The LLM-as-a-Judge evaluation shows consistently large effect sizes favoring machine-optimized documentation across all dimensions except Technical Accuracy. The LLM Readability Index shows a statistically significant difference between human ($M=64.03$) and machine ($M=78.72$) documentation, with a paired samples t-test yielding $t=-3.706$, $p=0.005$. The effect size for this comparison is Cohen's $d=-1.172$, indicating a large practical significance. The 95% confidence interval for the mean difference in LRI scores is [4.89, 24.49], which does not include zero, further confirming the statistical significance of the finding.

The paired t-test results for individual LLM evaluation dimensions confirm the pattern observed in the effect sizes. Conciseness shows the most significant difference ($t=-3.58$, $p=0.006$), followed by LLM-Friendliness ($t=-3.42$, $p=0.007$), Completeness ($t=-3.08$, $p=0.014$), and Clarity ($t=-3.01$, $p=0.015$). Technical Accuracy shows the least significant difference ($t=-0.78$, $p=0.455$), which is not statistically significant at the 0.05 level. This pattern suggests that the LLM evaluator perceives the most substantial improvements in dimensions related to information presentation and formatting, while technical accuracy remains consistently high across both documentation types.

The largest effects are observed for Conciseness ($d=-1.347$) and LLM-Friendliness ($d=-1.371$), suggesting that machine-optimized documentation is perceived as substantially more concise and better suited for LLM consumption. The Clarity ($d=-1.002$) and Completeness ($d=-0.964$) dimensions also show large effects, indicating that LLM evaluators perceive machine documentation as clearer and more complete.

Technical Accuracy shows the smallest effect ($d=-0.245$), suggesting that both human and machine documentation maintain similar levels of technical correctness. This is expected, as both versions describe the same technical systems and APIs.

Platform-level analysis reveals that 8 out of 10 platforms (80%) show higher LLM preference for machine-optimized documentation. Vercel is the notable exception, with human documentation scoring higher (54.00 vs. 42.00). This exception may be attributed to Vercel's machine documentation being particularly link-heavy with minimal prose content, reducing its perceived completeness and clarity. LangChain

TABLE IV: Platform-by-Platform LRI Comparison

Platform	Human LRI	Machine LRI	Delta	Direction
Stripe	55.00	80.00	+25.00	Machine
FastAPI	87.00	100.00	+13.00	Machine
Hono	73.33	87.00	+13.67	Machine
Cursor	75.00	100.00	+25.00	Machine
Supabase	56.00	82.00	+26.00	Machine
GitHub	55.00	81.25	+26.25	Machine
Vercel	54.00	42.00	-12.00	Human
Cloudflare	63.00	80.00	+17.00	Machine
React	72.00	85.00	+13.00	Machine
LangChain	50.00	50.00	0.00	Tie

Delta = Machine LRI - Human LRI. Positive values indicate machine documentation preference.

shows no difference (50.00 vs. 50.00), suggesting that both versions are equally optimized for LLM consumption.

FastAPI and Cursor achieve perfect LRI scores (100.00) for machine documentation, indicating optimal LLM-friendliness according to the evaluator. This may reflect the maturity of their llm.txt implementations and the completeness of their machine-optimized content.

The GitHub platform shows a substantial improvement from human (55.00) to machine (81.25) documentation, with a delta of +26.25 points. This is particularly notable given that GitHub's documentation is extremely extensive and covers a vast array of features. The machine-optimized version successfully condenses this information into a format that the LLM evaluator finds substantially more useful. Similarly, Stripe improves from 55.00 to 80.00 (+25.00), and Supabase from 56.00 to 82.00 (+26.00), indicating that payment and backend-as-a-service platforms benefit greatly from machine-optimized documentation.

Cloudflare's improvement from 63.00 to 80.00 (+17.00) and React's from 72.00 to 85.00 (+13.00) suggest that infrastructure and frontend frameworks also benefit from machine optimization, though the effect sizes are smaller than for developer platforms. This may reflect that these platforms have more technically sophisticated audiences, and their human documentation is already relatively well-optimized for technical readers.

The Hono platform shows an improvement from 73.33 to 87.00 (+13.67), indicating that even relatively newer frameworks can benefit from machine-optimized documentation. The fact that Hono's human documentation already scores relatively high (73.33) suggests that the framework's documentation team has invested in quality, but the machine-optimized version still provides additional value for LLM consumption.

The LangChain platform presents a unique case where both human and machine documentation achieve identical LRI scores of 50.00. This parity suggests that LangChain's human documentation is already well-optimized for LLM consumption, possibly due to its AI-native audience and the documentation team's awareness of LLM consumption patterns. The equal scores also indicate that the machine-optimized version does not provide additional value beyond the human-readable version, raising questions about the necessity of maintaining separate machine-optimized documentation when human documentation already achieves high LLM-friendliness. This case suggests that the readability-tokenization paradox may be less pronounced for platforms that naturally serve AI-literate audiences, and that the decision to implement llm.txt should be guided by the existing quality

of human documentation rather than applied uniformly across all platforms.

D. Correlation Analysis

The Pearson correlation analysis between traditional readability metrics and LLM evaluation scores reveals no significant correlations ($p > 0.05$) across all tested pairs. The highest correlation is observed between Flesch Reading Ease and LLM-Friendliness ($r=0.357$, $p=0.312$), but this does not reach statistical significance.

This lack of correlation provides critical evidence that traditional readability metrics and LLM evaluation scores measure fundamentally different constructs. While Flesch Reading Ease and Flesch-Kincaid Grade Level were designed to predict human reading comprehension based on sentence complexity and syllable density, LLM evaluation scores capture semantic quality, information density, and machine-specific optimization factors that traditional metrics cannot assess.

E. The Readability-Tokenization Paradox

The findings provide empirical evidence for what we term the “readability-tokenization paradox”: machine-optimized documentation that reduces token count and improves LLM performance simultaneously decreases traditional human readability metrics. This paradox is demonstrated through the inverse relationship between Flesch Reading Ease and LLM Readability Index:

- Human documentation: Higher FRE ($M=33.39$), Lower LRI ($M=64.03$)
- Machine documentation: Lower FRE ($M=-58.18$), Higher LRI ($M=78.72$)

This paradox has major implications for documentation strategy. Technical writers and documentation platform developers must balance the competing demands of human readability and machine optimization. The current trend toward llm.txt adoption suggests that the industry prioritizes machine optimization, potentially at the expense of human readability. However, this prioritization may not be optimal if the ultimate goal is to serve both human developers and AI assistants.

The paradox also raises questions about the evolving nature of technical communication. As LLMs mediate more of human access to technical information, the distinction between “human-readable” and “machine-readable” documentation becomes more blurred. When a developer asks an AI assistant about a library’s API, the assistant consumes the machine-optimized documentation but presents it in a human-friendly format. In this scenario, the LLM acts as a translator between machine-optimized source material and human-readable output, potentially rendering traditional human-readable documentation less critical for the end user.

F. Implications and Interpretation

The results suggest that the documentation landscape is undergoing a fundamental shift from human-centric to machine-optimized design. While traditional readability metrics remain valuable for assessing human comprehension, they are insufficient for evaluating documentation designed for dual human-machine consumption.

The large effect sizes for LLM evaluation dimensions (Cohen’s $d > -0.96$) indicate that machine-optimized documentation is not merely marginally better but substantially superior from an LLM perspective. This has practical implications for documentation teams: investing in llm.txt optimization can materially improve LLM performance and user experience when documentation is consumed through AI assistants. The effect sizes suggest that the improvement is practically significant, not just statistically significant, meaning that real-world users will notice meaningful differences in LLM response quality.

The Vercel exception provides a particularly instructive case. Vercel’s human documentation scored higher ($LRI=54.00$) than its machine documentation ($LRI=42.00$), suggesting that the machine-optimized version was not effectively implemented for that platform. This exception highlights that the readability-tokenization paradox is not universal; it depends on implementation quality and content characteristics. Platforms that invest in high-quality machine-optimized documentation reap the benefits of improved LLM performance, while those with poorly implemented machine documentation may actually degrade both human and machine readability.

The outliers in React and LangChain ($FRE=-430.09$ and -234.73 respectively) require further interpretation. These extreme negative values suggest that the machine-optimized documentation for these platforms contains extremely dense, technical content with minimal explanatory prose. React’s documentation, for instance, is heavily code-centric, with extensive API reference material that naturally scores poorly on traditional readability metrics. However, the LLM evaluation scores for these platforms (React: $LRI=85.00$, LangChain: $LRI=50.00$) suggest that LLMs can still process this dense content effectively, though LangChain’s machine documentation did not achieve the high LLM scores that other platforms did.

The lack of correlation between traditional metrics and LLM scores (all $p > 0.05$) provides the most theoretically important finding. This lack of correlation suggests that the two measurement systems are orthogonal: traditional metrics capture surface-level linguistic complexity (sentence length, syllable count), while LLM scores capture semantic and pragmatic quality (clarity, completeness, conciseness, LLM-friendliness). This orthogonality has important methodological implications: researchers studying documentation quality cannot rely solely on traditional metrics and must incorporate machine-centric evaluation methodologies.

The strongest correlation observed, though still not statistically significant, is between Token-to-Word Ratio and LLM-Friendliness ($r=0.564$, $p=0.090$). This near-significant correlation suggests that token efficiency, a purely computational metric, may have some relationship with LLM-perceived quality. However, the correlation is driven by a few platforms with extreme token ratios (React: 5.57, LangChain: 5.43), and when these outliers are removed, the correlation drops to $r=0.312$, $p=0.412$. This sensitivity to outliers indicates that token efficiency alone is not a reliable predictor of LLM-perceived quality, and that other factors (such as content structure, clarity, and completeness) play more important roles.

The correlation between Flesch Reading Ease and LLM-

Friendliness ($r=0.357$, $p=0.312$) shows a weak positive trend, suggesting that documentation that is easier for humans to read may also be somewhat more friendly to LLMs. However, this relationship is not strong enough to be statistically significant, and the direction of the correlation is opposite to what the readability-tokenization paradox would predict. This paradoxical finding may be explained by the fact that both human-readable and machine-optimized documentation can achieve high LLM-Friendliness scores if they are well-structured and clear, regardless of their traditional readability scores.

The correlation between token-to-word ratio and LLM scores ($r=0.564$, $p=0.090$) approaches statistical significance, suggesting that token efficiency may be one bridge between traditional metrics and LLM performance. While not conclusive at the 0.05 threshold, this trend indicates that documentation that uses fewer tokens per word (i.e., more common, less fragmented vocabulary) may be perceived more favorably by LLMs. This finding aligns with research on tokenization efficiency and suggests that vocabulary choice, not just sentence structure, influences machine readability.

G. Practical Implications for Documentation Teams

The findings provide actionable guidance for technical documentation teams. First, teams should invest in creating high-quality llm.txt files alongside their human-readable documentation. The large LLM evaluation improvements (mean LRI difference of +14.69 points) suggest that this investment yields measurable improvements in LLM performance. Second, teams should not abandon human-readable documentation entirely, as the Vercel exception demonstrates that poorly implemented machine documentation can degrade both human and machine readability.

Third, the lack of correlation between traditional metrics and LLM scores suggests that documentation teams should use both human and machine evaluation metrics to assess documentation quality. Relying solely on Flesch Reading Ease or Grade Level may miss critical aspects of machine readability, while relying solely on LLM evaluation may miss aspects of human comprehension. A dual-metric approach provides the most complete assessment.

Fourth, the extreme variance in traditional metrics for machine documentation ($SD=155.28$ for FRE) suggests that machine-optimized documentation is not uniform across platforms. Teams should aim for consistent quality in their machine-optimized documentation, avoiding the extreme density that produces extremely negative readability scores. The ideal machine-optimized documentation balances conciseness with sufficient explanatory context to maintain readability for both humans and machines.

H. Theoretical Implications for Technolinguistics

The readability-tokenization paradox identified in this study contributes to the emerging field of technolinguistics by demonstrating that the relationship between human and machine language processing is not simply one of alignment or opposition. The paradox suggests that optimization for one type of processor (human or machine) can degrade performance for the other, but the degradation is not symmetric. Machine-optimized documentation degrades traditional human

readability metrics markedly, but human-readable documentation does not necessarily degrade LLM performance (as evidenced by the Vercel exception, where human documentation scored higher on LLM evaluation).

This asymmetry suggests that the current generation of LLMs is more adaptable to diverse text formats than human readers are. LLMs can process dense, technical, poorly structured text and still extract meaningful information, while humans require clear structure, appropriate vocabulary, and sufficient context. This adaptability has implications for the future of technical communication: as LLMs improve, the pressure to optimize documentation for human readability may decrease, potentially leading to a bifurcation where documentation exists in two parallel forms (human-optimized and machine-optimized) rather than a single dual-purpose format.

The study also introduces the LLM Readability Index (LRI) as a standardized metric for machine readability assessment. While preliminary, this metric provides a framework for future research to compare documentation quality across platforms and evaluate the effectiveness of machine-optimized documentation standards. The LRI normalizes the 1-5 Likert scale to a 0-100 range, enabling direct comparison with traditional metrics and facilitating cross-platform analysis.

V. CONCLUSION

This study provides empirical evidence for the readability-tokenization paradox in technical documentation. Through analysis of ten platforms using both traditional readability metrics and a novel LLM-as-a-Judge methodology, we demonstrate that machine-optimized documentation (llm.txt) markedly improves LLM-perceived quality while generally decreasing traditional human readability metrics.

The key findings are:

- 1) **RQ1:** Machine-optimized documentation exhibits significantly different traditional readability metrics, with Flesch Reading Ease showing mixed results due to outlier sensitivity, but Flesch-Kincaid Grade Level consistently indicating higher reading difficulty (+10.2 grade levels).
- 2) **RQ2:** LLM-as-a-Judge scores strongly favor machine-optimized documentation ($p=0.005$), with large effect sizes across all evaluation dimensions (Cohen's d ranging from -0.964 to -1.371).
- 3) **RQ3:** Traditional readability metrics are insufficient for evaluating LLM-optimized documentation, as evidenced by the lack of significant correlation between traditional metrics and LLM scores (all $p > 0.05$).

These findings contribute to the emerging field of technolinguistics by quantifying the tension between human and machine optimization in documentation design. The LLM Readability Index (LRI) introduced in this study provides a standardized metric for future research comparing human and machine documentation quality.

The study's methodological contribution is equally important. The automated documentation auditing tool developed for this study provides a reproducible framework for large-scale documentation analysis. The tool's modular architecture, full test suite, and caching mechanisms make it suitable for longitudinal studies tracking documentation quality over time.

The integration of multiple evaluation methodologies (traditional metrics, LLM evaluation, statistical analysis) within a single pipeline demonstrates the feasibility of thorough documentation assessment at scale.

From a practical standpoint, the findings suggest that documentation teams should adopt a dual-format strategy: maintain human-readable HTML documentation for direct developer consumption while simultaneously investing in high-quality `llm.txt` files for LLM-mediated consumption. The marked improvement in LLM evaluation scores (mean LRI increase of +14.69 points) demonstrates that this investment yields measurable returns in terms of AI assistant performance. The Vercel exception, however, cautions that poorly implemented machine documentation can actually degrade both human and machine readability, emphasizing the importance of quality in machine-optimized documentation.

The theoretical implications of the readability-tokenization paradox extend beyond documentation to other forms of technical communication. As LLMs become increasingly central to information retrieval and processing, the tension between human and machine optimization will likely manifest in other domains: academic papers, legal documents, medical records, and scientific reports. The methodology developed in this study provides a template for investigating these tensions in other contexts, contributing to the broader field of human-machine communication.

A. Limitations

This study has several methodological limitations that should be acknowledged to ensure appropriate interpretation of the findings.

Construct Validity of Traditional Metrics on Non-Prose Content. The most important limitation concerns the construct validity of Flesch Reading Ease and Flesch-Kincaid Grade Level when applied to machine-optimized documentation. Analysis of the raw corpus reveals that several machine-optimized files (e.g., React, LangChain, Cloudflare) consist primarily of structured link directories rather than prose paragraphs. React’s machine-optimized file (`llm.txt`) contains 250 lines of which over 90% are bullet-pointed links with minimal descriptive text; LangChain’s machine file is an API endpoint listing with 959 entries, functionally equivalent to a directory index. Traditional readability metrics were developed for continuous prose and assume the presence of sentences, paragraphs, and natural language flow. When applied to link-heavy, structurally sparse content, these metrics produce extreme artifacts: React’s FRE score of -430.09 reflects not genuine unreadability but rather the absence of sentence-level content for the metric to evaluate. This means the traditional readability scores for machine-optimized documentation capture a mixture of two confounded factors: (1) genuine linguistic complexity, and (2) structural format differences between prose and directory-style content. The large standard deviation in machine FRE scores ($SD = 155.28$) is driven disproportionately by these structural artifacts rather than by variance in actual linguistic difficulty.

Content Heterogeneity and Implementation Quality. The ten platforms in the corpus exhibit substantial heterogeneity in how they implement `llm.txt` standards. At one extreme, FastAPI and Cursor provide complete, prose-rich

machine-optimized documentation that closely parallels the depth of their human-readable counterparts. At the other extreme, Vercel’s machine-optimized file is a 147,577-line link directory with minimal explanatory text, and the LLM evaluator assigned it a lower LRI score (42.00) than the human-readable version (54.00). This suggests that the “machine-optimized” label is not a uniform quality category but rather a spectrum of implementation maturity. The large variance in machine documentation quality (LRI range: 42.00 to 100.00; FRE range: -430.09 to 33.13) indicates that the comparison between “human” and “machine” documentation is, in practice, a comparison between relatively standardized human HTML and highly variable machine implementations. Future research should control for implementation quality by stratifying platforms based on documentation maturity or by experimentally manipulating documentation format while holding content constant.

Single-Model LLM Evaluation and Measurement Uncertainty. The LLM-as-a-Judge methodology employs a single evaluator (nvidia/nemotron-nano-12b-v2-v1, 12B parameters) operating at temperature 0.0 for deterministic outputs. While this model was selected for accessibility and reproducibility, its 12B parameter scale may limit its capacity to distinguish subtle differences in documentation quality, particularly for highly technical or structurally unconventional content. The LangChain result—where both human and machine documentation received identical scores (50.00) across all five evaluation dimensions—suggests that the evaluator may have defaulted to neutral ratings when confronted with unfamiliar or structurally complex API documentation. The near-ceiling scores for FastAPI and Cursor (100.00) similarly raise questions about score compression at the upper bound. Multi-model evaluation using larger models (e.g., Llama 3.1 70B, GPT-4) would test whether these patterns are robust across model scales, as recommended by Chen et al. [23] in their analysis of LLM judges for scientific programming.

Temporal Validity and Documentation Drift. Documentation is a dynamic artifact; platforms update their content continuously. The data collection for this study occurred during a single temporal window (early 2025), and the `llm.txt` standard itself was only formalized in September 2024. Platforms such as Vercel and Cloudflare may still be iterating on their machine-optimized documentation implementations, meaning the results reflect a snapshot of an evolving ecosystem rather than a stable equilibrium. The GitHub machine documentation, for instance, improved from -6.62 FRE to 81.25 LRI, suggesting a relatively mature implementation, while Vercel’s negative LRI delta suggests an immature one. Longitudinal replication would be necessary to disentangle platform-specific effects from genuine format-level effects.

Sampling and Generalizability Constraints. The purposive sample of ten platforms was selected based on `llm.txt` availability and domain diversity, not through random sampling. This limits statistical generalizability to the broader population of technical documentation. All platforms are English-language, and the domains represented (frontend frameworks, cloud infrastructure, AI tools, payment systems) exclude important categories such as enterprise software, hardware documentation, and scientific computing libraries.

The English-only restriction is particularly relevant given that Flesch-based metrics are calibrated for English syllable structure and may not generalize to morphologically complex languages.

Data Cleaning Pipeline Artifacts. The human-readable documentation underwent automated cleaning (Mozilla Readability extraction, code block removal, navigation stripping) before metric calculation. While this ensures fair comparison by removing HTML structural elements, the cleaning pipeline may not be equally effective across all platform types. Single-page applications (e.g., Cursor) required JavaScript rendering, which may introduce rendering artifacts or incomplete content extraction. The 500-character minimum content threshold, while necessary for meaningful metric calculation, may have excluded platforms with sparse documentation, potentially biasing the sample toward more complete documentation sites.

Correlation Analysis and Causal Inference. The correlation analysis between traditional metrics and LLM scores should be interpreted cautiously. The lack of significant correlation ($p > 0.05$) does not demonstrate that the constructs are independent; rather, it suggests that the sample size ($n = 10$) may lack sufficient power to detect moderate correlations. A post-hoc power analysis for Pearson correlation with $\alpha = 0.05$ and $n = 10$ reveals that the minimum detectable correlation is approximately $r = 0.63$, meaning correlations below this magnitude would not reach significance even if they exist in the population. The observational design also precludes causal claims about whether machine-optimized formatting directly causes lower traditional readability scores or whether both are driven by a third variable (e.g., platform technical complexity).

B. Future Work

Future research should expand the corpus to 50+ platforms, include non-English documentation, and employ multi-model LLM evaluation for consensus scoring. Comparative analysis with Markdown-formatted documentation and user studies examining actual human reading performance would provide additional validation. The development of hybrid metrics that simultaneously capture human and machine readability remains an important research direction.

Specifically, the following research directions are recommended: (1) Cross-platform validation: Replicate this study on a larger corpus of 50+ platforms to confirm the generalizability of the readability-tokenization paradox. (2) Multi-model evaluation: Employ 3-5 different LLM evaluators (e.g., GPT-4, Claude, Llama, Gemini) to establish consensus scores and assess model-specific biases. (3) Non-English documentation: Extend the methodology to Chinese, Japanese, and European languages to assess cross-linguistic applicability. (4) Longitudinal analysis: Track documentation quality over time as platforms evolve their llm.txt implementations. (5) User studies: Conduct controlled experiments with human developers to measure actual comprehension and task completion times using human vs. machine-optimized documentation. (6) Hybrid metrics: Develop composite metrics that combine traditional readability scores with LLM evaluation scores to create a unified documentation quality index. (7) Economic analysis: Quantify the cost-benefit tradeoffs of maintaining dual documentation formats, considering API costs, developer

time, and user satisfaction. (8) Integration studies: Investigate how llm.txt documentation affects the performance of downstream AI applications, such as code generation and technical support chatbots.

ACKNOWLEDGMENT

The author acknowledges the use of AI tools for statistical analysis, visualization, and literature review synthesis. The LLM-as-a-Judge evaluation was conducted using the nvidia/nemotron-nano-12b-v2-vl model via the OpenRouter API. This research was submitted to CODHES 2026 [24]. The author maintains full responsibility for research design, data interpretation, and conclusion formulation.

DATA AVAILABILITY STATEMENT

The data that support the findings of this study are openly available in the GitHub repository at <https://github.com/yehezkialgunawan/codhes-2026>, including raw scraped texts, statistical analysis scripts, and generated charts.

REFERENCES

- [1] J. Howard. (2024, September 3) The /llms.txt file: A proposal to standardise on using an /llms.txt file to provide information to help LLMs use a website at inference time. [Online]. Available: <https://llmstxt.org/>
- [2] AnswerDotAI. (2024) llms.txt. GitHub. [Online]. Available: <https://github.com/AnswerDotAI/llms-txt>
- [3] (2024) llmstxt.site. [Online]. Available: <https://llmstxt.site/>
- [4] (2024) directory.llmstxt.cloud. [Online]. Available: <https://directory.llmstxt.cloud/>
- [5] R. Flesch, "A new readability yardstick," *Journal of Applied Psychology*, vol. 32, no. 3, pp. 221–233, 1948.
- [6] J. P. Kincaid, R. P. Fishburne, R. L. Rogers, and B. S. Chissom, "Derivation of new readability formulas (automated readability index, fog count and flesch reading ease formula) for navy enlisted personnel," Naval Air Station Memphis, Research Branch Report 8-75, 1975. [Online]. Available: <https://apps.dtic.mil/sti/pdfs/ADA006655.pdf>
- [7] J. C. Redish, "Readability formulas have even more limitations than Klare discusses," *ACM Journal of Computer Documentation*, vol. 24, no. 3, pp. 132–137, 2000.
- [8] G. McClure, "Readability formulas: Useful or useless? (an interview with J. Peter Kincaid)," *IEEE Transactions on Professional Communication*, vol. 30, no. 1, pp. 12–15, 1987.
- [9] J. P. Kincaid, R. Braby, and J. Mears, "Electronic authoring and delivery of technical information," *Journal of Instructional Development*, vol. 11, no. 2, pp. 8–13, 1988.
- [10] R. Sennrich, B. Haddow, and A. Birch, "Neural machine translation of rare words with subword units," in *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (ACL 2016)*, Berlin, Germany, 2016, pp. 1715–1725.
- [11] OpenAI. (2023) tiktoken: A fast BPE tokeniser for use with OpenAI's models. GitHub. [Online]. Available: <https://github.com/openai/tiktoken>
- [12] Z. Ankner, A. Bertsch, G. Neubig, and M. R. Gormley, "The ":" tokenization," in *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
- [13] S. Ding, A. Renduchintala, and K. Duh, "A call for prudent choice of subword merge operations in neural machine translation," in *Proceedings of MT Summit 2019*, 2019. [Online]. Available: <https://aclanthology.org/W19-6608/>
- [14] N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang, "Lost in the middle: How language models use long contexts," *Transactions of the Association for Computational Linguistics*, vol. 11, 2023.
- [15] H. Jiang, Q. Wu, C.-Y. Lin, Y. Yang, and L. Qiu, "LLMLingua: Compressing prompts for accelerated inference of large language models," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, 2023, pp. 18 020–18 036.
- [16] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, and I. Stoica, "Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena," in *Proceedings of the Neural Information Processing Systems (NeurIPS) 2023 Datasets and Benchmarks Track*, 2023.

- [17] K. Horikawa, K. Horikawa, Y. Kashiwa, H. Uwano, and H. Iida, "Do AI agents really improve code readability?" in *Proceedings of the 23rd International Conference on Mining Software Repositories (MSR '26)*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.13723>
- [18] I. R. d. S. Simoes and E. Venson, "Measuring how changes in code readability attributes affect code quality evaluation by large language models," in *Proceedings of the XXXIX Brazilian Symposium on Software Engineering*, 2025.
- [19] C. Diggs, M. Doyle, A. Madan, S. Scott, E. Escamilla, J. Zimmer, N. Nekoo, P. Ursino, M. Bartholf, Z. Robin, A. Patel, C. Glasz, W. Macke, P. Kirk, J. Phillips, A. Sridharan, D. Wendt, S. Rosen, N. Naik, J. F. Brunelle, and S. Thaker, "Leveraging LLMs for legacy code modernization: Challenges and opportunities for LLM-Generated documentation," in *Proceedings of the LLM4Code Workshop at ICSE 2025*, 2024. [Online]. Available: <https://arxiv.org/abs/2411.14971>
- [20] Textstat. (2023) textstat: Calculate statistical features from text. Python package. [Online]. Available: <https://github.com/shivam5992/textstat>
- [21] Mozilla. (2024) Readability: A standalone version of the readability library used for Firefox reader view. Mozilla. [Online]. Available: <https://github.com/mozilla/readability>
- [22] (2024) Openrouter: A unified interface for LLMs. [Online]. Available: <https://openrouter.ai/>
- [23] A. Chen, C. Wong, B. Sharif, and A. Peruma, "Exploring code comprehension in scientific programming: Preliminary insights from research scientists," in *Proceedings of the 33rd IEEE/ACM International Conference on Program Comprehension (ICPC 2025)*, 2025. [Online]. Available: <https://arxiv.org/abs/2501.10037>
- [24] (2026) CODHES 2026: International conference on digital humanities and social sciences. Binus University. [Online]. Available: <https://humanities.binus.ac.id/codhes/>